

# ECE 47300 Assignment 5

Your Name:Tina Xu

Prepare the packages we will use.

```
In [22]: import time
from typing import List, Dict

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torchvision
import torchvision.models as models
import torchvision.transforms as transforms

import matplotlib.pyplot as plt
```

## Exercise 0: Train your model on GPU (0 points)

For some tasks in this assignment, it can take a long time if you run it on CPU. For example, based on our test on Exercise 3 Task 4, it will take roughly 2 hours to train the full model for 1 epoch on CPU. Hence, we highly recommend you try to train your model on GPU.

To do so, first you need to enable GPU on Colab (this will restart the runtime). Click **Runtime** -> **Change runtime type** and select the **Hardware accelerator** there. You can then run the following code to see if the GPU is correctly initialized and available.

**Note:** If you would like to avoid GPU overages on Colab, we would suggest writing and debugging your code before switching on the GPU runtime. Otherwise, the time you spent debugging code will likely count against your GPU usage. Once you have the code running, you can switch on the GPU runtime and train the model much faster.

```
In [23]: print(f'Can I can use GPU now? -- {torch.cuda.is_available()}')

Can I can use GPU now? -- True
```

## You must manually move your model and data to the GPU (and sometimes back to the cpu)

After setting the GPU up on colab, then you should put your **model** and **data** to GPU. We give a simple example below. You can use `to` function for this task. See `torch.Tensor.to` to move a tensor to the GPU (probably your mini-batch of data in each iteration) or `torch.nn.Module.to` to move your NN model to GPU (assuming you create subclass `torch.nn.Module`). Note that `to()` of tensor returns a NEW tensor while `to` of a NN model will apply this in-place. To be safe, the

best semantics are `obj = obj.to(device)`. For printing, you will need to move a tensor back to the CPU via the `cpu()` function.

Once the model and input data are on the GPU, everything else can be done the same. This is the beauty of PyTorch GPU acceleration. None of the other code needs to be altered.

To summarize, you need to 1) enable GPU acceleration in Colab, 2) put the model on the GPU, and 3) put the input data (i.e., the batch of samples) onto the GPU using `to()` after it is loaded by the data loaders (usually you only put one batch of data on the GPU at a time).

```
In [24]: rand_tensor = torch.rand(5,2)
simple_model = nn.Sequential(nn.Linear(2,10), nn.ReLU(), nn.Linear(10,1))
print(f'input is on {rand_tensor.device}')
print(f'model parameters are on {[param.device for param in simple_model.parameters()]}')
print(f'output is on {simple_model(rand_tensor).device}')

# device = torch.device('cuda')
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# ----- <Your code> -----
# Move rand_tensor and model onto the GPU device
rand_tensor = rand_tensor.to(device)
simple_model = simple_model.to(device)

# ----- <End your code> -----
print(f'input is on {rand_tensor.device}')
print(f'model parameters are on {[param.device for param in simple_model.parameters()]}')
print(f'output is on {simple_model(rand_tensor).device}')
```

```
input is on cpu
model parameters are on [device(type='cpu'), device(type='cpu'), device(type='cpu'),
device(type='cpu')]
output is on cpu
input is on cuda:0
model parameters are on [device(type='cuda', index=0), device(type='cuda', index=0),
device(type='cuda', index=0), device(type='cuda', index=0)]
output is on cuda:0
```

## Exercise 1: Why use a CNN rather than only fully connected layers? (40 points)

In this exercise, you will build two models for the **MNIST** dataset: one uses only fully connected layers and another uses a standard CNN layout (convolution layers everywhere except the last layer is fully connected layer). Note, you will need to use cross entropy loss as your objective function. The two models should be built with roughly the same accuracy performance, your task is to compare the number of network parameters (a huge number of parameters can affect training/testing time, memory requirements, overfitting, etc.).

### Task 1: Prepare train and test function

We will create our train and test procedure in these two functions. The train function should apply one epoch of training. The functions inputs should take everything we need for training

and testing and return some logs.

### Arguments requirement:

- For the `train` function, it takes the `model`, `loss_fn`, `optimizer`, `train_loader`, and `epoch` as arguments.
  - `model`: the classifier, or deep neural network, should be an instance of `nn.Module`.
  - `loss_fn`: the loss function instance. For example, `nn.CrossEntropy()`, or `nn.L1Loss()`, etc.
  - `optimizer`: should be an instance of `torch.optim.Optimizer`. For example, it could be `optim.SGD()` or `optim.Adam()`, etc.
  - `train_loader`: should be an instance of `torch.utils.data.DataLoader`.
  - `epoch`: the current number of epoch. Only used for log printing.(default: 1.)
- For the `test` function, it takes all the inputs above except for the optimizer (and it takes a test loader instead of a train loader).

### Log requirement:

Here are some further requirements:

- In the `train` function, print the log 8-10 times per epoch. The print statement should be:
 

```
print(f'Epoch {epoch}:  
[{batch_idx*len(images)}/{len(train_loader.dataset)}] Loss:  
{loss.item():.3f}')
```
- In the `test` function, print the log after the testing. The print statement is:
 

```
print(f"Test result on epoch {epoch}: total sample: {total_num}, Avg  
loss: {test_stat['loss']:.3f}, Acc: {100*test_stat['accuracy']:.3f}%")
```

### Return requirement

- The `train` function should return a list, which the element is the loss per batch, i.e., one loss value for every batch.
- The `test` function should return a dictionary with three keys: "loss", "accuracy", and "prediction". The values are the average loss of all the testset, average accuracy of all the test dataset, and the prediction of all test dataset.

### Other requirement:

- In the `train` function, the model should be updated in-place, i.e., do not copy the model inside `train` function.

```
In [25]: def train(model: nn.Module,
              loss_fn: nn.modules.loss._Loss,
              optimizer: torch.optim.Optimizer,
              train_loader: torch.utils.data.DataLoader,
              epoch: int=0)-> List:
    # ----- <Your code> -----
    model.train()
```

```

train_loss = []
train_counter = []
for batch_idx, (images, label) in enumerate(train_loader):
    images, label = images.to(device), label.to(device)
    label = label.squeeze()
    optimizer.zero_grad()
    output = model(images)
    loss = loss_fn(output, label)
    loss.backward()
    optimizer.step()

    #if batch_idx % 10 == 0: # We record our output every 10 batches
    train_loss.append(loss.item())
    if batch_idx % 100 == 0:
        print(f'Epoch {epoch}: [{batch_idx*len(images)}/{len(train_loader.dataset)}] L
# ----- <End Your code> -----
assert len(train_loss) == len(train_loader)
return train_loss

def test(model: nn.Module,
        loss_fn: nn.modules.loss._Loss,
        test_loader: torch.utils.data.DataLoader,
        epoch: int=0)-> Dict:
    # ----- <Your code> -----
    total_num = len(test_dataset)
    model.eval()
    test_loss = 0
    correct = 0
    pred_list = []
    with torch.no_grad():
        for images, targets in test_loader:
            images, targets = images.to(device), targets.to(device)
            output = model(images)
            test_loss += loss_fn(output, targets)
            pred = output.data.max(1, keepdim=True)[1]
            pred_list.extend(pred.view(-1).cpu().numpy())
            correct += pred.eq(targets.data.view_as(pred)).sum()

    test_loss /= len(test_loader.dataset)
    accuracy = correct / len(test_loader.dataset)
    test_stat = {
        'loss': test_loss,
        'accuracy': accuracy,
        'prediction': torch.tensor(pred_list)
    }

    print(f"Test result on epoch {epoch}: total sample: {total_num}, Avg loss: {test_s
# ----- <End Your code> -----
# dictionary should include loss, accuracy and prediction
assert "loss" and "accuracy" and "prediction" in test_stat.keys()
# "prediction" value should be a 1D tensor
assert len(test_stat["prediction"]) == len(test_loader.dataset)
assert isinstance(test_stat["prediction"], torch.Tensor)
return test_stat

```

**Task 2: Following the structure used in the instructions, you should create**

## One network named `OurFC` which should consist with only fully connected layers

- You need to add one `nn.Linear(*, 256)` where one of the dimension is `256` and decide how many other layers and hidden dimensions you want in your network (apart from this).
- Your final accuracy on the test dataset should lie roughly around 97% ( $\pm 2\%$ )
- There is no need to make the neural network unnecessarily complex, your total training time should no longer than 3 mins

## Another network named `OurCNN` which applies a standard CNN structure

- You should have one `nn.Conv2d(*, *, kernel_size=5)` convolutional layer with `kernel_size=5`, and again, you should decide how many layers and channels you want for each layer.
- Your final accuracy on the test dataset should lie roughly around 97% ( $\pm 2\%$ )
- A standard CNN structure can be composed as **[Conv2d, MaxPooling, ReLU] x num\_conv\_layers + FC x num\_fc\_layers**
- Train and test your network on MNIST data as in the instructions.
- Notice You can always use the `train` and `test` function you write throughout this assignment.
- The code below will also print out the number of parameters for both neural networks to allow comparison.
- (You can use multiple cells if helpful but make sure to run all of them to receive credit.)

```
In [26]: # Download MNIST and transformation
# ----- <Your code> -----
transform = torchvision.transforms.Compose([torchvision.transforms.ToTensor(),
                                           torchvision.transforms.Normalize((0.1307,), (0.3081,))])

train_dataset = torchvision.datasets.MNIST('data', train=True, download=True, transform=transform)
test_dataset = torchvision.datasets.MNIST('data', train=False, download=True, transform=transform)

print(train_dataset)
# ----- <End Your code> -----
```

```
Dataset MNIST
  Number of datapoints: 60000
  Root location: data
  Split: Train
  StandardTransform
Transform: Compose(
  ToTensor()
  Normalize(mean=(0.1307,), std=(0.3081,))
)
```

```
In [27]: # Build OurFC class and OurCNN class.
# ----- <Your code> -----
class OurFC(nn.Module):
    def __init__(self):
        super(OurFC, self).__init__()
```

```

self.fc1 = nn.Linear(28 * 28, 256)
self.fc2 = nn.Linear(256, 128)
self.fc3 = nn.Linear(128, 64)
self.fc4 = nn.Linear(64, 10)

def forward(self, x):
    x = x.view(-1, 28 * 28)
    x = F.relu(self.fc1(x))
    x = F.relu(self.fc2(x))
    x = F.relu(self.fc3(x))
    x = self.fc4(x)
    return x

class OurCNN(nn.Module):
    def __init__(self):
        super(OurCNN, self).__init__()
        self.conv = nn.Conv2d(1, 3, kernel_size=5)
        self.conv2 = nn.Conv2d(3, 6, kernel_size=3)
        self.fc = nn.Linear(432, 10)

    def forward(self, x):
        x = self.conv(x)          # x now has shape (batchsize x 3 x 24 x 24)
        x = F.relu(F.max_pool2d(x,2)) # x now has shape (batchsize x 3 x 12 x 12)
        x = x.view(-1, 432)      # x now has shape (batchsize x 432)
        x = F.relu(self.fc(x))   # x has shape (batchsize x 10)
        return F.log_softmax(x, -1)
# ----- <End Your code> -----

```

```

In [28]: # Let's first train the FC model. Below are there common hyperparameters.
criterion = nn.CrossEntropyLoss()

start = time.time()
max_epoch = 4
# ----- <Your code> -----
batch_size_train, batch_size_test = 64, 1000
train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size_train,
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size_test, sh

max_epoch = 3
classifier = OurFC()
classifier = classifier.to(device)
optimizer = optim.SGD(classifier.parameters(), lr=0.01, momentum=0.8)
train_loss = []
test_losses = []

for epoch in range(1, max_epoch+1):
    train(classifier, criterion, optimizer, train_loader, epoch)
    test(classifier, criterion, test_loader, epoch)

# ----- <End Your code> -----
end = time.time()
print(f'Finished Training after {end-start} s ')

```

```

Epoch 1: [0/60000] Loss: 2.3225107192993164
Epoch 1: [6400/60000] Loss: 0.9476959705352783
Epoch 1: [12800/60000] Loss: 0.6406949162483215
Epoch 1: [19200/60000] Loss: 0.35390621423721313
Epoch 1: [25600/60000] Loss: 0.19334815442562103
Epoch 1: [32000/60000] Loss: 0.38602402806282043
Epoch 1: [38400/60000] Loss: 0.29433339834213257
Epoch 1: [44800/60000] Loss: 0.19536735117435455
Epoch 1: [51200/60000] Loss: 0.24671359360218048
Epoch 1: [57600/60000] Loss: 0.3474026918411255
Test result on epoch 1: total sample: 10000, Avg loss: 0.000, Acc: 94.600%
Epoch 2: [0/60000] Loss: 0.06289239972829819
Epoch 2: [6400/60000] Loss: 0.12653708457946777
Epoch 2: [12800/60000] Loss: 0.13868169486522675
Epoch 2: [19200/60000] Loss: 0.09734631329774857
Epoch 2: [25600/60000] Loss: 0.16340899467468262
Epoch 2: [32000/60000] Loss: 0.22193831205368042
Epoch 2: [38400/60000] Loss: 0.04736936092376709
Epoch 2: [44800/60000] Loss: 0.10074757784605026
Epoch 2: [51200/60000] Loss: 0.15586017072200775
Epoch 2: [57600/60000] Loss: 0.06606201827526093
Test result on epoch 2: total sample: 10000, Avg loss: 0.000, Acc: 96.290%
Epoch 3: [0/60000] Loss: 0.1526002436876297
Epoch 3: [6400/60000] Loss: 0.1196659579873085
Epoch 3: [12800/60000] Loss: 0.06038293242454529
Epoch 3: [19200/60000] Loss: 0.06527547538280487
Epoch 3: [25600/60000] Loss: 0.05436798185110092
Epoch 3: [32000/60000] Loss: 0.26116129755973816
Epoch 3: [38400/60000] Loss: 0.1038118377327919
Epoch 3: [44800/60000] Loss: 0.1607113480567932
Epoch 3: [51200/60000] Loss: 0.1048082560300827
Epoch 3: [57600/60000] Loss: 0.03551391884684563
Test result on epoch 3: total sample: 10000, Avg loss: 0.000, Acc: 97.120%
Finished Training after 46.12486934661865 s

```

```

In [29]: criterion = nn.CrossEntropyLoss()

# Let's then train the OurCNN model.
start = time.time()
# ----- <Your code> -----
max_epoch = 3
classifier = OurCNN()
classifier = classifier.to(device)
optimizer = optim.SGD(classifier.parameters(), lr=0.01, momentum=0.8)
train_loss = []
test_loss = []

for epoch in range(1, max_epoch+1):
    train(classifier, criterion, optimizer, train_loader, epoch)
    test(classifier, criterion, test_loader, epoch)

# ----- <End Your code> -----
end = time.time()
print(f'Finished Training after {end-start} s ')

```

```

Epoch 1: [0/60000] Loss: 2.3162405490875244
Epoch 1: [6400/60000] Loss: 0.8069671988487244
Epoch 1: [12800/60000] Loss: 0.8876321315765381
Epoch 1: [19200/60000] Loss: 0.6190297603607178
Epoch 1: [25600/60000] Loss: 0.5085657835006714
Epoch 1: [32000/60000] Loss: 0.48061224818229675
Epoch 1: [38400/60000] Loss: 0.26104214787483215
Epoch 1: [44800/60000] Loss: 0.22581329941749573
Epoch 1: [51200/60000] Loss: 0.07480457425117493
Epoch 1: [57600/60000] Loss: 0.1471724659204483
Test result on epoch 1: total sample: 10000, Avg loss: 0.000, Acc: 96.150%
Epoch 2: [0/60000] Loss: 0.24099670350551605
Epoch 2: [6400/60000] Loss: 0.12015841901302338
Epoch 2: [12800/60000] Loss: 0.14905524253845215
Epoch 2: [19200/60000] Loss: 0.12678638100624084
Epoch 2: [25600/60000] Loss: 0.05848908796906471
Epoch 2: [32000/60000] Loss: 0.04194008186459541
Epoch 2: [38400/60000] Loss: 0.1255013644695282
Epoch 2: [44800/60000] Loss: 0.1321403831243515
Epoch 2: [51200/60000] Loss: 0.22423624992370605
Epoch 2: [57600/60000] Loss: 0.18203610181808472
Test result on epoch 2: total sample: 10000, Avg loss: 0.000, Acc: 96.610%
Epoch 3: [0/60000] Loss: 0.07076451927423477
Epoch 3: [6400/60000] Loss: 0.017751364037394524
Epoch 3: [12800/60000] Loss: 0.11676153540611267
Epoch 3: [19200/60000] Loss: 0.06384273618459702
Epoch 3: [25600/60000] Loss: 0.07712694257497787
Epoch 3: [32000/60000] Loss: 0.03436969593167305
Epoch 3: [38400/60000] Loss: 0.06634313613176346
Epoch 3: [44800/60000] Loss: 0.0653982013463974
Epoch 3: [51200/60000] Loss: 0.07489781081676483
Epoch 3: [57600/60000] Loss: 0.24033835530281067
Test result on epoch 3: total sample: 10000, Avg loss: 0.000, Acc: 97.430%
Finished Training after 45.245373487472534 s

```

```

In [30]: ourfc = OurFC()
total_params = sum(p.numel() for p in ourfc.parameters())
print(f'OurFC has a total of {total_params} parameters')

ourcnn = OurCNN()
total_params = sum(p.numel() for p in ourcnn.parameters())
print(f'OurCNN has a total of {total_params} parameters')

```

```

OurFC has a total of 242762 parameters
OurCNN has a total of 4576 parameters

```

Questions (0 points, just for understanding): Which one has more parameters? Which one is likely to have less computational cost when deployed? Which one took longer to train?



## Exercise 2: Train classifier on CIFAR-10 data. (30 points)

Now, let's move our dataset to color images. CIFAR-10 dataset is another widely used dataset. Here all images have colors, i.e. each image has 3 color channels instead of only one channel in MNIST. You need to pay more attention to the dimension of the data as it passes through the layers of your network.

### Task 1: Create data loaders

- Load CIFAR10 train and test data with appropriate composite transform where the normalize transform should be `transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))`.
- Set up a `train_loader` and `test_loader` for the CIFAR-10 data with a batch size of 9 similar to the instructions.
- The code below will plot a 3 x 3 subplot of images including their labels. (do not modify)

```
In [31]: classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# Create the appropriate transform, load/download CIFAR10 train and test datasets with
# ----- <Your code> -----
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
    download=True, transform=transform)
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
    download=True, transform=transform)

train_loader = torch.utils.data.DataLoader(trainset, batch_size=9,
    shuffle=True, num_workers=2)
test_loader = torch.utils.data.DataLoader(testset, batch_size=9,
    shuffle=False, num_workers=2)

# ----- <End Your code> -----

# Code to display images
batch_idx, (images, targets) = next(enumerate(train_loader)) #fix!!!!
fig, ax = plt.subplots(3,3,figsize = (9,9))
for i in range(3):
    for j in range(3):
        image = images[i*3+j].permute(1,2,0)
        image = image/2 + 0.5
        ax[i,j].imshow(image)
        ax[i,j].set_axis_off()
        ax[i,j].set_title(f'{classes[targets[i*3+j]]}')
fig.show()
```

Files already downloaded and verified  
Files already downloaded and verified

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```



## Task 2: Create CNN and train it

Set up a convolutional neural network and have your data trained on it. You have to decide all the details in your network, overall your neural network should meet the following standards to receive full credit:

- You should not use more than three convolutional layers and three fully connected layers
- Accuracy on the test dataset should be **above** 50%

```
In [32]: # Create CNN network.
# ----- <Your code> -----
class Net(nn.Module):#This class is from Lecture demo
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x))) # (N, 6, 14, 14)
        x = self.pool(F.relu(self.conv2(x))) # (N, 16, 5, 5)
        x = x.view(-1, 16 * 5 * 5) # (N, 400)
        x = F.relu(self.fc1(x)) # (N, 120)
        x = F.relu(self.fc2(x)) # (N, 84)
        x = self.fc3(x) # (N, 10)
        return x
net = Net()
net = net.to(device)

# ----- <End Your code> -----
```

```
In [33]: # Train your neural network here.
start = time.time()
max_epoch = 4
# ----- <Your code> -----
net = net.to(device) # This function on Lecture demo
for epoch in range(max_epoch):
    running_loss = 0.0
    for i, data in enumerate(train_loader, 0):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer = optim.SGD(net.parameters(), lr=0.01, momentum=0.8)
        optimizer.zero_grad()
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    if i % 2000 == 1999:
        print('[%d, %5d] loss: %.3f' %(epoch + 1, i + 1, running_loss / 2000))
        running_loss = 0.0

# ----- <End Your code> -----
output = test(net, criterion, test_loader, epoch)
```

```
end = time.time()
print(f'Finished Training after {end-start} s ')
```

```
[1, 2000] loss: 2.127
[1, 4000] loss: 1.767
[2, 2000] loss: 1.511
[2, 4000] loss: 1.453
[3, 2000] loss: 1.337
[3, 4000] loss: 1.322
[4, 2000] loss: 1.242
[4, 4000] loss: 1.234
```

Test result on epoch 3: total sample: 10000, Avg loss: 0.142, Acc: 55.090%

Finished Training after 78.66084146499634 s

## Task 3: Plot misclassified test images

Plot some misclassified images in your test dataset:

- select five images that are **misclassified** for `class_id` in `{1,3,5,7,9}` by your neural network, one image each (i.e., the true label is `class_id` but the predicted label is not `class_id`).
- label each images with true label and predicted label
- use `detach().cpu()` when plotting images if the image is in gpu

```
In [34]: total_images = 5
predictions = output['prediction']
targets = torch.tensor(testset.targets)
# ----- <Your code> -----
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
import matplotlib.pyplot as plt
import numpy as np
misclassified_images = {}
misclassified_preds = {}
misclassified_targets = {}

net.eval()
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        for idx, (image, label) in enumerate(zip(predicted, labels)):
            if image != label and label.item() in [1, 3, 5, 7, 9] and label.item() not in [0, 2, 4, 6, 8]:
                misclassified_images[label.item()] = images[idx]
                misclassified_preds[label.item()] = image.item()
                misclassified_targets[label.item()] = label.item()
            if len(misclassified_images) >= 5:
                break
        if len(misclassified_images) >= 5:
            break

fig, axes = plt.subplots(1, 5, figsize=(15, 3))
for ax, class_id in zip(axes, [1, 3, 5, 7, 9]):
    img = misclassified_images[class_id].cpu().numpy().transpose((1, 2, 0))
    img = np.clip(img, 0, 1)
    ax.imshow(img)
    ax.set_title(f'True: {classes[misclassified_targets[class_id]]}\nPred: {classes[misclassified_preds[class_id]]}')
```

```
ax.axis('off')

plt.show()

# ----- <End Your code> -----
```



Questions (0 points): Are the mis-classified images also misleading to human eyes?

## Exercise 3: Transfer Learning (30 points)

In practice, people won't train an entire CNN from scratch, because it is relatively rare to have a dataset of sufficient size (or sufficient computational power). Instead, it is common to pretrain a CNN on a very large dataset and then use the CNN either as an initialization or a fixed feature extractor for the task of interest.

In this task, you will learn how to use a pretrained CNN for CIFAR-10 classification.

### Task1: Load pretrained model

`torchvision.models` (<https://pytorch.org/vision/stable/models.html>) contains definitions of models for addressing different tasks, including: image classification, pixelwise semantic segmentation, object detection, instance segmentation, person keypoint detection and video classification.

First, you should load the **pretrained** ResNet-18 that has already been trained on [ImageNet](#) using `torchvision.models`. If you are interested in more details about Resnet-18, read this paper <https://arxiv.org/pdf/1512.03385.pdf>.

```
In [35]: resnet18 = models.resnet18(pretrained=True)
         resnet18 = resnet18.to(device)
```

### Task2: Create data loaders for CIFAR-10

Then you need to create a modified dataset and dataloader for CIFAR-10. Importantly, the model you load has been trained on **ImageNet** and it expects inputs as mini-batches of 3-channel RGB images of shape (3 x H x W), where H and W are expected to be **at least** 224. So you need to preprocess the CIFAR-10 data to make sure it has a height and width of 224. Thus, you should add a transform when loading the CIFAR10 dataset (see

`torchvision.transforms.Resize` ). This should be added appropriately to the `transform` you created in a previous task.

```
In [54]: # Create your dataloader here
# ----- <Your code> -----
from torchvision.transforms import v2

transforms = v2.Compose([
    v2.RandomResizedCrop(size=(224, 224), antialias=True),
    v2.RandomHorizontalFlip(p=0.5),
    v2.PILToTensor(),
    v2.ToDtype(torch.float32, scale=True),
    v2.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]),
])
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transforms)
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transforms)
train_loader = torch.utils.data.DataLoader(trainset, batch_size=9, shuffle=True, num_workers=4)
test_loader = torch.utils.data.DataLoader(testset, batch_size=9, shuffle=False, num_workers=4)

# ----- <End Your code> -----
```

Files already downloaded and verified

Files already downloaded and verified

### Task3: Classify test data on pretrained model

[link text](#) Use the model you load to classify the **test** CIFAR-10 data and print out the test accuracy.

Don't be surprised if the accuracy is bad!

```
In [37]: # ----- <Your code> -----
test(resnet18, criterion, test_loader, epoch)
# ----- <End Your code> -----
```

Test result on epoch 3: total sample: 10000, Avg loss: 1.348, Acc: 0.030%

```
Out[37]: {'loss': tensor(1.3482, device='cuda:0'),
          'accuracy': tensor(0.0003, device='cuda:0'),
          'prediction': tensor([653, 567, 147, ..., 383, 392, 268])}
```

### Task 4: Fine-tune (i.e., update) the pretrained model for CIFAR-10

Now try to improve the test accuracy. We offer several possible solutions:

(1) You can try to directly continue to train the model you load with the CIFAR-10 training data.

(2) For efficiency, you can try to freeze part of the parameters of the loaded models. For example, you can first freeze all parameters by

```
for param in model.parameters():
    param.requires_grad = False
```

and then unfreeze the last few layers by setting `somelayer.requires_grad=True` .

You are also welcome to try any other approach you can think of.

**Note:** You must print out the test accuracy and to get full credits, the test accuracy should be at least **80%**.

```
In [41]: # Directly train the whole model.
start = time.time()
#----- <Your code> -----
max_epoch = 4
optimizer = optim.SGD(filter(lambda p: p.requires_grad, resnet18.parameters()), lr=0.01)
for epoch in range(1, max_epoch+1):
    train(resnet18, criterion, optimizer, train_loader, epoch)
# ----- <End Your code> -----
test(resnet18, criterion, test_loader, epoch)
end = time.time()
print(f'Finished Training after {end-start} s ')
```

Epoch 1: [0/50000] Loss: 0.9849302768707275  
Epoch 1: [900/50000] Loss: 0.6891946196556091  
Epoch 1: [1800/50000] Loss: 0.547964334487915  
Epoch 1: [2700/50000] Loss: 0.8184277415275574  
Epoch 1: [3600/50000] Loss: 0.5266135334968567  
Epoch 1: [4500/50000] Loss: 0.8284181356430054  
Epoch 1: [5400/50000] Loss: 0.9807007312774658  
Epoch 1: [6300/50000] Loss: 0.5416517853736877  
Epoch 1: [7200/50000] Loss: 0.4476836025714874  
Epoch 1: [8100/50000] Loss: 0.4538692831993103  
Epoch 1: [9000/50000] Loss: 0.66401606798172  
Epoch 1: [9900/50000] Loss: 0.9981821775436401  
Epoch 1: [10800/50000] Loss: 0.6260478496551514  
Epoch 1: [11700/50000] Loss: 0.6677668690681458  
Epoch 1: [12600/50000] Loss: 0.7393826246261597  
Epoch 1: [13500/50000] Loss: 0.7403638362884521  
Epoch 1: [14400/50000] Loss: 0.5619577169418335  
Epoch 1: [15300/50000] Loss: 0.8158794045448303  
Epoch 1: [16200/50000] Loss: 1.4424934387207031  
Epoch 1: [17100/50000] Loss: 0.45140504837036133  
Epoch 1: [18000/50000] Loss: 0.6159303188323975  
Epoch 1: [18900/50000] Loss: 0.28817957639694214  
Epoch 1: [19800/50000] Loss: 1.1235251426696777  
Epoch 1: [20700/50000] Loss: 0.23907096683979034  
Epoch 1: [21600/50000] Loss: 0.8489562273025513  
Epoch 1: [22500/50000] Loss: 1.2779484987258911  
Epoch 1: [23400/50000] Loss: 0.6469547748565674  
Epoch 1: [24300/50000] Loss: 1.697003960609436  
Epoch 1: [25200/50000] Loss: 0.5200656652450562  
Epoch 1: [26100/50000] Loss: 0.19058328866958618  
Epoch 1: [27000/50000] Loss: 0.3811015784740448  
Epoch 1: [27900/50000] Loss: 0.7860597372055054  
Epoch 1: [28800/50000] Loss: 0.3740847110748291  
Epoch 1: [29700/50000] Loss: 0.40207087993621826  
Epoch 1: [30600/50000] Loss: 0.36634576320648193  
Epoch 1: [31500/50000] Loss: 0.16625672578811646  
Epoch 1: [32400/50000] Loss: 1.1789175271987915  
Epoch 1: [33300/50000] Loss: 0.50594162940979  
Epoch 1: [34200/50000] Loss: 1.1102519035339355  
Epoch 1: [35100/50000] Loss: 0.2945215702056885  
Epoch 1: [36000/50000] Loss: 0.6855146884918213  
Epoch 1: [36900/50000] Loss: 0.7727756500244141  
Epoch 1: [37800/50000] Loss: 0.35939958691596985  
Epoch 1: [38700/50000] Loss: 0.4412614107131958  
Epoch 1: [39600/50000] Loss: 0.5531730055809021  
Epoch 1: [40500/50000] Loss: 0.5961208343505859  
Epoch 1: [41400/50000] Loss: 0.3412613272666931  
Epoch 1: [42300/50000] Loss: 0.48015096783638  
Epoch 1: [43200/50000] Loss: 0.7083131074905396  
Epoch 1: [44100/50000] Loss: 1.476594090461731  
Epoch 1: [45000/50000] Loss: 0.3266851305961609  
Epoch 1: [45900/50000] Loss: 0.6296382546424866  
Epoch 1: [46800/50000] Loss: 0.24130858480930328  
Epoch 1: [47700/50000] Loss: 1.169346570968628  
Epoch 1: [48600/50000] Loss: 0.34805166721343994  
Epoch 1: [49500/50000] Loss: 0.6282298564910889  
Epoch 2: [0/50000] Loss: 0.20323872566223145  
Epoch 2: [900/50000] Loss: 0.22967730462551117  
Epoch 2: [1800/50000] Loss: 0.38426321744918823  
Epoch 2: [2700/50000] Loss: 0.8651399612426758



Epoch 2: [3600/50000] Loss: 0.4233931005001068  
Epoch 2: [4500/50000] Loss: 0.5420306324958801  
Epoch 2: [5400/50000] Loss: 1.1341089010238647  
Epoch 2: [6300/50000] Loss: 0.21644242107868195  
Epoch 2: [7200/50000] Loss: 0.8338583707809448  
Epoch 2: [8100/50000] Loss: 0.8227366209030151  
Epoch 2: [9000/50000] Loss: 0.7115365862846375  
Epoch 2: [9900/50000] Loss: 0.5526877641677856  
Epoch 2: [10800/50000] Loss: 0.7473070025444031  
Epoch 2: [11700/50000] Loss: 1.6697274446487427  
Epoch 2: [12600/50000] Loss: 0.865280032157898  
Epoch 2: [13500/50000] Loss: 0.49589619040489197  
Epoch 2: [14400/50000] Loss: 0.1453690379858017  
Epoch 2: [15300/50000] Loss: 0.43032947182655334  
Epoch 2: [16200/50000] Loss: 0.5002106428146362  
Epoch 2: [17100/50000] Loss: 0.5027185678482056  
Epoch 2: [18000/50000] Loss: 0.8088065981864929  
Epoch 2: [18900/50000] Loss: 0.5112119317054749  
Epoch 2: [19800/50000] Loss: 0.5504079461097717  
Epoch 2: [20700/50000] Loss: 0.7669381499290466  
Epoch 2: [21600/50000] Loss: 0.29870861768722534  
Epoch 2: [22500/50000] Loss: 0.6886396408081055  
Epoch 2: [23400/50000] Loss: 0.6568844318389893  
Epoch 2: [24300/50000] Loss: 0.3270705044269562  
Epoch 2: [25200/50000] Loss: 0.9025945663452148  
Epoch 2: [26100/50000] Loss: 0.05027858540415764  
Epoch 2: [27000/50000] Loss: 0.44731923937797546  
Epoch 2: [27900/50000] Loss: 0.4595891535282135  
Epoch 2: [28800/50000] Loss: 0.41101211309432983  
Epoch 2: [29700/50000] Loss: 0.34626272320747375  
Epoch 2: [30600/50000] Loss: 0.5789348483085632  
Epoch 2: [31500/50000] Loss: 0.5101833343505859  
Epoch 2: [32400/50000] Loss: 0.2846188545227051  
Epoch 2: [33300/50000] Loss: 0.6277625560760498  
Epoch 2: [34200/50000] Loss: 0.878325879573822  
Epoch 2: [35100/50000] Loss: 0.4823019802570343  
Epoch 2: [36000/50000] Loss: 1.2615007162094116  
Epoch 2: [36900/50000] Loss: 0.21819984912872314  
Epoch 2: [37800/50000] Loss: 0.615347146987915  
Epoch 2: [38700/50000] Loss: 0.40865838527679443  
Epoch 2: [39600/50000] Loss: 0.3973275423049927  
Epoch 2: [40500/50000] Loss: 0.8778776526451111  
Epoch 2: [41400/50000] Loss: 0.3870818614959717  
Epoch 2: [42300/50000] Loss: 0.031020402908325195  
Epoch 2: [43200/50000] Loss: 0.4159247875213623  
Epoch 2: [44100/50000] Loss: 0.6430480480194092  
Epoch 2: [45000/50000] Loss: 0.578813374042511  
Epoch 2: [45900/50000] Loss: 0.03151388093829155  
Epoch 2: [46800/50000] Loss: 1.8487226963043213  
Epoch 2: [47700/50000] Loss: 0.44276708364486694  
Epoch 2: [48600/50000] Loss: 1.1476013660430908  
Epoch 2: [49500/50000] Loss: 0.6635501384735107  
Epoch 3: [0/50000] Loss: 0.8140881061553955  
Epoch 3: [900/50000] Loss: 1.3710519075393677  
Epoch 3: [1800/50000] Loss: 0.5944159030914307  
Epoch 3: [2700/50000] Loss: 1.125868320465088  
Epoch 3: [3600/50000] Loss: 0.150893896818161  
Epoch 3: [4500/50000] Loss: 0.3372439742088318  
Epoch 3: [5400/50000] Loss: 0.31924259662628174  
Epoch 3: [6300/50000] Loss: 0.3049107789993286

Epoch 3: [7200/50000] Loss: 0.36053428053855896  
Epoch 3: [8100/50000] Loss: 0.2757761776447296  
Epoch 3: [9000/50000] Loss: 1.0924931764602661  
Epoch 3: [9900/50000] Loss: 0.5736520886421204  
Epoch 3: [10800/50000] Loss: 1.0072927474975586  
Epoch 3: [11700/50000] Loss: 0.6616535782814026  
Epoch 3: [12600/50000] Loss: 0.5944721102714539  
Epoch 3: [13500/50000] Loss: 0.5050820112228394  
Epoch 3: [14400/50000] Loss: 0.46904468536376953  
Epoch 3: [15300/50000] Loss: 0.5426397323608398  
Epoch 3: [16200/50000] Loss: 0.7912757992744446  
Epoch 3: [17100/50000] Loss: 0.47037798166275024  
Epoch 3: [18000/50000] Loss: 1.0176444053649902  
Epoch 3: [18900/50000] Loss: 1.0751872062683105  
Epoch 3: [19800/50000] Loss: 1.4605355262756348  
Epoch 3: [20700/50000] Loss: 1.0882437229156494  
Epoch 3: [21600/50000] Loss: 0.3398069441318512  
Epoch 3: [22500/50000] Loss: 0.6517194509506226  
Epoch 3: [23400/50000] Loss: 0.15191291272640228  
Epoch 3: [24300/50000] Loss: 0.6915101408958435  
Epoch 3: [25200/50000] Loss: 1.0269641876220703  
Epoch 3: [26100/50000] Loss: 0.09979817271232605  
Epoch 3: [27000/50000] Loss: 0.8432350158691406  
Epoch 3: [27900/50000] Loss: 0.3602488040924072  
Epoch 3: [28800/50000] Loss: 1.089316725730896  
Epoch 3: [29700/50000] Loss: 1.0595890283584595  
Epoch 3: [30600/50000] Loss: 0.1678881049156189  
Epoch 3: [31500/50000] Loss: 0.33670851588249207  
Epoch 3: [32400/50000] Loss: 1.0444612503051758  
Epoch 3: [33300/50000] Loss: 0.30274391174316406  
Epoch 3: [34200/50000] Loss: 1.1024484634399414  
Epoch 3: [35100/50000] Loss: 0.1848711520433426  
Epoch 3: [36000/50000] Loss: 0.39340299367904663  
Epoch 3: [36900/50000] Loss: 0.8191666603088379  
Epoch 3: [37800/50000] Loss: 0.6369274258613586  
Epoch 3: [38700/50000] Loss: 0.6730743050575256  
Epoch 3: [39600/50000] Loss: 0.3020130395889282  
Epoch 3: [40500/50000] Loss: 0.7070930004119873  
Epoch 3: [41400/50000] Loss: 1.0670198202133179  
Epoch 3: [42300/50000] Loss: 0.5216058492660522  
Epoch 3: [43200/50000] Loss: 0.23438532650470734  
Epoch 3: [44100/50000] Loss: 0.6135627627372742  
Epoch 3: [45000/50000] Loss: 0.349598228931427  
Epoch 3: [45900/50000] Loss: 0.26562485098838806  
Epoch 3: [46800/50000] Loss: 0.8983584642410278  
Epoch 3: [47700/50000] Loss: 1.1109256744384766  
Epoch 3: [48600/50000] Loss: 0.1894942969083786  
Epoch 3: [49500/50000] Loss: 0.36178258061408997  
Epoch 4: [0/50000] Loss: 0.3306840658187866  
Epoch 4: [900/50000] Loss: 0.30540451407432556  
Epoch 4: [1800/50000] Loss: 0.8343163728713989  
Epoch 4: [2700/50000] Loss: 0.30153611302375793  
Epoch 4: [3600/50000] Loss: 0.9777965545654297  
Epoch 4: [4500/50000] Loss: 0.18115361034870148  
Epoch 4: [5400/50000] Loss: 1.1445307731628418  
Epoch 4: [6300/50000] Loss: 0.28653040528297424  
Epoch 4: [7200/50000] Loss: 0.5045759677886963  
Epoch 4: [8100/50000] Loss: 0.6300418972969055  
Epoch 4: [9000/50000] Loss: 1.1178699731826782  
Epoch 4: [9900/50000] Loss: 0.6410528421401978

```

Epoch 4: [10800/50000] Loss: 0.5219218730926514
Epoch 4: [11700/50000] Loss: 0.3320443332195282
Epoch 4: [12600/50000] Loss: 0.23285868763923645
Epoch 4: [13500/50000] Loss: 1.387982726097107
Epoch 4: [14400/50000] Loss: 0.8041008114814758
Epoch 4: [15300/50000] Loss: 0.19343164563179016
Epoch 4: [16200/50000] Loss: 0.2519959807395935
Epoch 4: [17100/50000] Loss: 0.5802217721939087
Epoch 4: [18000/50000] Loss: 0.5105597376823425
Epoch 4: [18900/50000] Loss: 0.7861250638961792
Epoch 4: [19800/50000] Loss: 1.0933032035827637
Epoch 4: [20700/50000] Loss: 0.467133104801178
Epoch 4: [21600/50000] Loss: 0.8308685421943665
Epoch 4: [22500/50000] Loss: 0.35176241397857666
Epoch 4: [23400/50000] Loss: 0.3832857310771942
Epoch 4: [24300/50000] Loss: 0.38329702615737915
Epoch 4: [25200/50000] Loss: 1.3424688577651978
Epoch 4: [26100/50000] Loss: 1.9836993217468262
Epoch 4: [27000/50000] Loss: 0.4520920217037201
Epoch 4: [27900/50000] Loss: 1.003834843635559
Epoch 4: [28800/50000] Loss: 0.6368727684020996
Epoch 4: [29700/50000] Loss: 1.8800686597824097
Epoch 4: [30600/50000] Loss: 0.15716761350631714
Epoch 4: [31500/50000] Loss: 0.597868025302887
Epoch 4: [32400/50000] Loss: 0.9818665981292725
Epoch 4: [33300/50000] Loss: 0.9539517164230347
Epoch 4: [34200/50000] Loss: 1.2056217193603516
Epoch 4: [35100/50000] Loss: 1.1459869146347046
Epoch 4: [36000/50000] Loss: 0.2410164624452591
Epoch 4: [36900/50000] Loss: 1.340208888053894
Epoch 4: [37800/50000] Loss: 0.7258832454681396
Epoch 4: [38700/50000] Loss: 0.28668299317359924
Epoch 4: [39600/50000] Loss: 0.9338753819465637
Epoch 4: [40500/50000] Loss: 1.2484253644943237
Epoch 4: [41400/50000] Loss: 0.27167269587516785
Epoch 4: [42300/50000] Loss: 0.6388657689094543
Epoch 4: [43200/50000] Loss: 0.2011478692293167
Epoch 4: [44100/50000] Loss: 1.4204493761062622
Epoch 4: [45000/50000] Loss: 1.1570276021957397
Epoch 4: [45900/50000] Loss: 0.37552720308303833
Epoch 4: [46800/50000] Loss: 0.1384652554988861
Epoch 4: [47700/50000] Loss: 0.16978037357330322
Epoch 4: [48600/50000] Loss: 0.35698598623275757
Epoch 4: [49500/50000] Loss: 0.8229844570159912
Test result on epoch 4: total sample: 10000, Avg loss: 0.061, Acc: 81.720%
Finished Training after 294.74332642555237 s

```

```

In [57]: # Load another resnet18 instance, only unfreeze the outer layers.
# ----- <Your code> -----
criterion = nn.CrossEntropyLoss()
epoch = 0
for param in resnet18.parameters():
    param.requires_grad = False
num = resnet18.fc.in_features
resnet18.fc = nn.Linear(num, 10)
resnet18.layer4.requires_grad = True
#resnet18.layer3.requires_grad = True
resnet18.fc.requires_grad = True
# ----- <End Your code> -----

```

```

In [58]: # Train the model!!
start = time.time()
# ----- <Your code> -----
resnet18 = resnet18.to(device)
num_epochs = 1
for epoch in range(num_epochs):
    resnet18.train()
    running_loss = 0.0
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        optimizer = optim.SGD(filter(lambda p: p.requires_grad, resnet18.parameters()))
        optimizer.zero_grad()
        outputs = resnet18(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        running_loss += loss.item()
    print(f'Epoch {epoch+1}, Loss: {running_loss/len(train_loader)}')
# ----- <End Your code> -----
test(resnet18, criterion, test_loader)
end = time.time()
print(f'Finished Training after {end-start} s ')

```

Epoch 1, Loss: 1.0273716919381597

Test result on epoch 0: total sample: 10000, Avg loss: 0.072, Acc: 80.340%

Finished Training after 85.05671691894531 s